

Homework 2

Tommaso Pasini & Valentina Pyatkin

(pasini | pyatkin)@di.uniroma1.it

Supervised POS Tagger with LSTM

The Tasks

- 1. Implement a sequence POS tagger with a LSTM.**
- 2. Extends your model to cope with other languages. (Extra Points)**

What we provide:

- We will provide you a folder that has the following structure:
 - homework2/
 - slides.pdf
 - src/
 - homework2.py
 - data/
 - en-ud-train.conllu
 - en-ud-test.conllu
 - en-ud-dev.conllu

data/ - Universal Dependencies TreeBank

- Universal Dependencies (<http://universaldependencies.org/>) is a project that is developing cross-linguistically consistent treebank annotation for many languages.
- The tag set contains 17 POS.
- All the languages have the same tag set.
- This enables multi-lingual POS-tagging.
- CoNLL format is used. (<http://universaldependencies.org/format.html>)

Input Data - The CoNLL Format

- CoNLL format has one token per line.
- Each line is divided in 10 column tab separated:
 - 1 Word ID: starts from 1 for each sentence and may be a range for multiwords
 - 2 FORM: the word form or the punctuation symbol.
 - 3 LEMMA: the lemma of the word form.
 - 4 U-POSTAG: universal part of speech.
 - 5 X-POSTAG: language specific pos tag (_ if not defined).
 - 6 FEATS: list of morphological features from universal feature inventory.
 - 7 HEAD: Word ID of the head of the word.
 - 8 DEPREL: universal dependency relation.
 - 9 DEPS: enhanced dependency graph.
 - 10 MISC: any other annotation.

Input Data - The CoNLL Format

- CoNLL format has one token per line.
- Each line is divided in 10 column tab separated:
 - 1 Word ID: starts from 1 for each sentence and may be a range for multiwords
 - 2 FORM: the word form or the punctuation symbol.
 - 3 LEMMA: the lemma of the word form.
 - 4 U-POSTAG: universal part of speech.
 - 5 X-POSTAG: language specific pos tag (_ if not defined).
 - 6 FEATS: list of morphological features from universal feature inventory.
 - 7 HEAD: Word ID of the head of the word.
 - 8 DEPREL: universal dependency relation.
 - 9 DEPS: enhanced dependency graph.
 - 10 MISC: any other annotation.

Input Data - The CoNLL Format

- UPOS

- XPOS

- FORM

- LEMMA

- WORD ID

1	>	Al	Al	PROPN	NNP	Number=Sing	0	root	_	SpaceAfter=No
2	>	-	-	PUNCT	HYPH	_	1	punct	_	SpaceAfter=No
3	>	Zaman	Zaman	PROPN	NNP	Number=Sing	1	flat	_	-
4	>	:	:	PUNCT	:	_	1	punct	_	-
5	>	American	american	ADJ	JJ	Degree=Pos	6	amod	_	-
6	>	forces	force	NOUN	NNS	Number=Plur	7	nsubj	_	-
7	>	killed	kill	VERB	VBD	Mood=Ind Tense=Past VerbForm=Fin	1	parataxis	_	-
8	>	Shaikh	Shaikh	PROPN	NNP	Number=Sing	7	obj	_	-
9	>	Abdullah	Abdullah	PROPN	NNP	Number=Sing	8	flat	_	-
10	>	al	al	PROPN	NNP	Number=Sing	8	flat	_	SpaceAfter=No
11	>	-	-	PUNCT	HYPH	_	8	punct	_	SpaceAfter=No
12	>	Ani	Ani	PROPN	NNP	Number=Sing	8	flat	_	SpaceAfter=No
13	>	,	,	PUNCT	,	_	8	punct	_	-
14	>	the	the	DET	DT	Definite=Def PronType=Art	15	det	_	-
15	>	preacher	preacher	NOUN	NN	Number=Sing	8	appos	_	-
16	>	at	at	ADP	IN	_	18	case	_	-
17	>	the	the	DET	DT	Definite=Def PronType=Art	18	det	_	-
18	>	mosque	mosque	NOUN	NN	Number=Sing	7	obl	_	-
19	>	in	in	ADP	IN	_	21	case	_	-
20	>	the	the	DET	DT	Definite=Def PronType=Art	21	det	_	-
21	>	town	town	NOUN	NN	Number=Sing	18	nmod	_	-
22	>	of	of	ADP	IN	_	23	case	_	-
23	>	Qaim	Qaim	PROPN	NNP	Number=Sing	21	nmod	_	SpaceAfter=No
24	>	,	,	PUNCT	,	_	21	punct	_	-
25	>	near	near	ADP	IN	_	28	case	_	-
26	>	the	the	DET	DT	Definite=Def PronType=Art	28	det	_	-
27	>	Syrian	syrian	ADJ	JJ	Degree=Pos	28	amod	_	-
28	>	border	border	NOUN	NN	Number=Sing	21	nmod	_	SpaceAfter=No
29	>	.	.	PUNCT	.	_	1	punct	_	-

Input Data - Tag Set

- Universal dependencies define 17 different POS tags:

(<http://universaldependencies.org/u/pos/index.html>)

- ADJ: adjective
- ADP: adposition
- ADV: adverb
- AUX: auxiliary
- CONJ: coordinating conjunction
- DET: determiner
- INTJ: interjection
- NOUN: noun
- NUM: numeral
- PART: particle
- PRON: pronoun
- PROPN: proper noun
- PUNCT: punctuation
- SCONJ: subordinating conjunction
- SYM: symbol
- VERB: verb
- X: other

Input Data - Split

- We will provide you the data.
- Data will be splitted in:
 - Training data: 229.672 tokens.
 - Development data: 29.152 tokens.
 - Test data: 29.250 tokens.
- Do not merge them!

homework2.py

- homework2.py define 3 abstract classes:
 - AbstractPOSTaggerTrainer
 - AbstractPOSTaggerTester
 - AbstractLSTMPOSTagger
- Homework2.py also define a Test class that implements a test method.
 - Run the homework2.py in order to test your implementation:
`python homework2.py [--no-train] <homework_dir> <model_path>`
- **You need to pass the test in order to be evaluated on this homework.**
- homework2.py also contains an utility class (ModelIO) which implements methods to save and load a model.

What and how to code:

- You have to **implement** the 3 python **abstract classes** provided in homework2.py:
 - POSTaggerTrainer
 - POSTaggerTester
 - LSTMPOSTagger
- Each method in each class must be implemented.
- The implementation must be tested with the test() method in the test class Test

AbstractPOSTaggerTrainer:

- AbstractPOSTaggerTrainer represents a class to train a model.
- Has 1 abstract method `train(training_path)` which takes as input the path to the training data and outputs a Sequential model.

```
class AbstractPOSTaggerTrainer:
    __metaclass__ = ABCMeta

    @abstractmethod
    def train(self, training_path):
        """
        Train the keras model from the training data.

        :param training_path: the path to training file
        :return: the keras model of type Sequential
        """
        pass
```

AbstractPOSTaggerTester:

- AbstractPOSTaggerTester represents a class to **test a trained model**.
- Has 1 abstract method `test(model, test_file_path)` which takes as input a trained sequential model and the path to the gold standard data. It **outputs a dictionary** containing **precision, recall, coverage** and **f1**.

```
class AbstractPOSTaggerTester:
    __metaclass__ = ABCMeta

    @abstractmethod
    def test(self, model, test_file_path):
        """
        Test the input model against the gold standard.

        :param model: a Sequential model that has to be tested.
        :param test_file_path: a path to the gold standard file.

        :return: a dictionary that has as keys 'precision', 'recall',
        'coverage' and 'f1' and as associated value their respective values.

        Additional info:
        - Precision has to be computed as the number of correctly predicted
          pos tag over the number of predicted pos tags.
        - Recall has to be computed as the number of correctly predicted
          pos tag over the number of items in the gold standard
        - Coverage has to be computed as the number of predicted pos tag over
          the number of items in the gold standard
        - F1 has to be computed as the armonic mean between precision
          and recall ( $2 * P * R / (P + R)$ )
        """
        pass
```

AbstractLSTMPOSTagger:

- AbstractLSTMPOSTagger is an abstract pos tagger that can predict the pos tags given a tokenized sentence.
- Has 1 abstract method predict(sentence) which given a list of tokens outputs a list of same length with the pos tags for each word.

```
class AbstractLSTMPOSTagger:
    __metaclass__ = ABCMeta

    def __init__(self, model):
        self._model = model

    def get_model(self):
        return self._model

    @abstractmethod
    def predict(self, sentence):
        """
        predict the pos tags for each token in the sentence.
        :param sentence: a list of tokens.
        :return: a list of pos tags (one for each input token).
        """
        pass
```

ModelIO:

- ModelIO it's an utility class and implements two methods:
 - `save(model, path)` that save the model on the specified path.
 - `load(path)` that load a keras model from the path
- You should use this class in order to ensure that you save and load your model correctly.
- This class will be also used to test that your implementations return the correct objects.

```
class ModelIO:
    @staticmethod
    def save(model, output_path):
        """
        Save the model to in the file pointed by the output_path variable

        :param model: the trained Sequential model
        :param output_path: the path to the file on which the model have to
                           be saved
        :return: no return value is required
        """
        model.save(output_path)

    @staticmethod
    def load(model_file_path):
        """
        Load a sequential model saved in the file pointed by model_file_path

        :param model_file_path: the path to the file that has to be loaded
        :return: a sequential model loaded from the file
        """
        import keras
        return keras.models.load_model(model_file_path)
```


The Test class:

- Test class implements a simple type checking test method to ensure that the classes you have implemented return the correct objects.
- You have to run successfully this test in order to ensure that you correctly implemented the homework.
- If the test fail, your homework will not be evaluated.

```
class Test:
    def __init__(self, training_path, model_path, gold_standard_path):
        self._training_path = training_path
        self._model_path = model_path
        self._gold_standard_path = gold_standard_path

    def test(self, lstm_trainer_implementation, lstm_tester_implementation, no_train):
        if no_train:
            model = ModelIO.load(self._model_path)
            print 'TEST 0\t\tNO-TRAIN'
        else:
            model = lstm_trainer_implementation.train(self._training_path)
            assert type(model) == Sequential
            print 'TEST 0\t\tPASSED'
            ModelIO.save(model, self._model_path)
            model = ModelIO.load(self._model_path)

        assert type(model) == Sequential

        print 'TEST 1\t\tPASSED'

        results = lstm_tester_implementation.test(model, self._gold_standard_path)
        assert type(results) == dict
        assert 'precision' in results.keys()
        assert 'recall' in results.keys()
        assert 'coverage' in results.keys()
        assert 'f1' in results.keys()

        print 'TEST 2\t\tPASSED'

        postagger = LSTMPOSTagger(model)
        test_sentence = ['this', 'is', 'an', 'easy', 'test']
        prediction = postagger.predict(test_sentence)

        assert len(test_sentence) == len(prediction)

        print 'TEST 3\t\tPASSED'

        return results
```

How to run the test

- python homework2.py
/path/to/model.keras
/path/to/your/submission_folder/
- **Be careful!!!** If /path/to/model.keras exist and the *--no-train* option is not specified then it will be overridden.
- The
/path/to/your/submission_folder
should be structured as specified in the previous slide, especially, it should contains the src/ folder with your implementation and the data/ folder

```
if __name__ == '__main__':  
    """  
    Main to run the test of the homework 2.  
    Use the parameter --no-train in order to skip the training of the model  
    WARNING!!! the program will not check if model_path already exist and will overwrite it if it  
    is the case  
    """  
    if len(sys.argv) < 3:  
        print 'usage: python', sys.argv[0], '[--no-train] model_path, homework_dir'  
        sys.exit(-1)  
    model_index = 1  
    homework_dir_index = 2  
    no_train = False  
    if '--no-train' in sys.argv:  
        model_index = 2  
        homework_dir_index = 3  
        no_train = True  
  
    model_output_path = sys.argv[model_index]  
    homework_dir = sys.argv[homework_dir_index]  
  
    src_dir = homework_dir + 'src/'  
  
    print ''  
    print 'model output:', model_output_path  
    print 'homework dir:', homework_dir  
    print 'src dir:', src_dir  
    print ''  
  
    # dynamic import of modules  
    sys.path.append(src_dir)  
    from POSTaggerTester import POSTaggerTester  
    from LSTMPOSTagger import LSTMPOSTagger  
    from POSTaggerTrainer import POSTaggerTrainer  
  
    ## get files  
    training_data = homework_dir + '/data/en-ud-train.conllu'  
    test_data = homework_dir + '/data/en-ud-test.conllu'  
  
    test = Test(training_data, model_output_path, test_data)  
    trainer = POSTaggerTrainer()  
    tester = POSTaggerTester()  
  
    name = ''  
    if homework_dir.endswith('/'):   
        name = homework_dir[:-1]  
    else:  
        name = homework_dir  
  
    name = name[name.rfind("/"):]  
  
    try:  
        results = test.test(trainer, tester, no_train=no_train)  
        print results  
        print name, "PASSED"  
    except Exception as e:  
        print name, "FAILED"  
        raise traceback.print_exc(e)
```

What to submit

- You have to submit a folder structured as follow:
 - cognome_matricola_homework_2/
 - homework2.py
 - report.pdf
 - src/
 - POSTaggerTrainer.py
 - POSTaggerTester.py
 - LSTMPOSTagger.py
 - data/
 - en-ud-train.conllu
 - en-ud-test.conllu
 - en-ud-dev.conllu
 - output/
 - pos_tagged_sentences.txt
 - results.txt
 - model.keras

What to submit - src

- src/ has to contain your source code. Specifically it has to contain **one file for each implemented class** and the contained **class has to be named with the file name**.
- **POSTaggerTrainer** and **POSTaggerTester** must have a **default empty constructor**.
- **LSTMPostTagger** has to implement a constructor that takes as input a Sequential model.
- You can include additional files auxiliaries file.

What to submit - output

- **pos_tagged_sentences.txt:**

Must contain the output of LSTMPOSTagger#predict called on all test instances (From the test set). Has to be formatted as follow:

 sentence_1

 pos-tagged sentence_1

 ...

- **results.txt:**

Must contain a line for each measure (precision, recall, f1, coverage) with its score.

- **model.keras:** your trained model

The report

- In the root folder you have to include:
 - the **homework2.py** library that we provide.
 - the **report.pdf** in pdf format.
- The report must be of maximum 2 pages and has to follow the following format:

Abstract: a brief task description.

Model description: description of your NN model, what layers you used, how many layers, activation function ecc.

Optimization: description of how and which parameters you tuned; how and which word-vector representations you used.

Experiments: description of your test and discussion of the results. Create a confusion matrix (12 x 12) of the POStags in which each cell (i,j) contains the number of times the model predicted the i-th POStag but should had predicted the j-th POStag. (A heat map of the confusion matrix would be very appreciated). Discuss interesting findings from the confusion matrix, using examples.

Extra points

- Universal dependencies enables multilingual POS tagging.
- Trainig sets are available for other languages at <https://lindat.mff.cuni.cz/repository/xmlui/handle/11234/1-1983>
- We will provide word-embeddings for other languages.

Extra points - The Task

- Merge English and Italian universal training set.
- Download Italian pre-trained word embeddings:
http://hlt.isti.cnr.it/wordembeddings/skipgram_wiki_window10_size300_neg-samples10.tar.gz
- Use the same code you wrote to train your model with the new training set in Italian and English.
- Merge dev and test set of the two languages.
- Tune and Test the multilingual model on the new datasets.


Extra points - What to submit

- Add multilingual/ folder to the root of your submission.
- Multilingual has to be structured as follow:
 - multilingual/
 - data/
 - ud-merged-train.conll
 - ud-merged-test.conll
 - ud-merged-dev.conll
 - output/
 - pos_tagged_sentences.txt
 - results.txt
 - multilingual_model.keras

How-To

How-To: transforming your labels into y vectors

- Reminder from practical session 4 and 5:
 - Do a **label-to-index mapping**
 - Do a **index-to-label mapping**
 - Represent your label as a **one-hot-encoding** using the label-to-index mapping
 - Retrieve the label from the y-vector given the index-to-label mapping
 - You can use zeros and ones, or booleans (like in the example)



```
[False False False False False False False False False False False False
False False False False False False False False False False False False
False False False False False False False False False False False False
False False False False False False False False True False False False
False False False False False False False False False False False]
```

How-To: Getting your input matrix

- Convert your words into appropriate vector representations that can be used as input to your NN (look at slides from practical sessions 4+5)
- Reminder from the last two practical sessions:
 - represent each word as a vector
 - add the word-vectors into another vector to represent the context of the temporal sequence you are looking at
 - add the temporal sequences to the final matrix that you will use as input for the network
 - ! your input is three-dimensional !

```
[[[False False False ..., False False False]
 [False False False ..., False False False]
 [False False False ..., False False False]
 ...,
 [False False False ..., False False False]
 [False False False ..., False False False]
 [False False False ..., False False False]]

[[[False False False ..., False False False]
 [False False False ..., False False False]
 [False False False ..., False False False]
 ...,
 [False False False ..., False False False]
 [False False False ..., False False False]
 [False False False ..., False False False]]

[[[False False False ..., False False False]
 [ True False False ..., False False False]
 [ True False False ..., False False False]
 ...,
 [False False False ..., False False False]
 [False False False ..., False False False]
 [False False False ..., False False False]]


...,
```

How-To: Word-Vector Representations

- There are many ways to represent your input words as vectors.
- It is up to **you** to choose an appropriate representation.
- The following slides will give you some ideas on what kind of vectors you can work with.

Word-Vector Representations: One-hot encodings

- Reminder: dimensionality of the length of the set of your predefined vocabulary
- Binary or boolean
- Cons: high dimensionality, sparse



```
[False False False False False False False False False False False False
False False False False False False False False False False False False
False False False False False False False False False False False False
False False False False False False False False True False False False
False False False False False False False False False False False]
```

Word-Vector Representations: GloVe

- Use pretrained GloVe vectors: <https://nlp.stanford.edu/projects/glove/>
- Distributional representation
- Word to word co-occurrence statistics in a corpus
- Cons: stopwords are not included... -> you have to solve this problem yourself:
 - a solution:
 - have a special "unknown-words"/UNK vector
- Pros: captures semantics of words
 - related words are closer in the (distributional) semantic space

Word-Vector Representations: Word2Vec

- Use pretrained Word2Vec vectors:
<https://github.com/mmhaltz/word2vec-GoogleNews-vectors>
- Neural model: Skip-Gram or CBOW with either hierarchical softmax or negative sampling
- Cons: stopwords are not included... -> you have to solve this problem yourself:
 - a solution:
 - have a special "unknown-words"/UNK vector
- Pros: captures semantics of words
 - embeddings
 - reduced dimensionality

Word-Vector Representations: Build-your-own

- Ideas:
 - Build your own vectors by taking **distributional counts**, over a fixed number of content words and context, and weighting them with **tfidf** (term frequency - inverse document frequency)
 - You could also work with other dimensionality reduction methods like **PCA** and **LSA**

The Neural Network Layers

- **Input layer**

- takes your word vector representations
 - Input dimension: your vector dimensions (depend on the representation you chose)
 - Output dimension: same (it's a linear layer)

- **Embedding layer:**

- Input: one hot representation of the word.
- Output: latent (embedded) vector with the chosen size.

- **LSTM layer**

- Long-Short-Term-Memory Network
 - you can define your own hidden layer size

- **Output layer**

- does a softmax over the hidden layer from the LSTM, in order to predict the most probable PoS
 - you get a matrix with *batchsize* \times *y-vector-size* dimensions, e.g. every vector in your matrix is a probability distribution for the PoS tags for a single word.

Our Tips

- Start by transforming your input into vector representation(s)
- Start with a simple network and see how it performs. Then make it more complex.
- Keep in mind that training NNs might take
—some time (start soon!)

The Deadline

- Sunday, 7th of May
- <http://robertonavigli.com/nlp2017/>